

TAKE-HOME ASSIGNMENT

Creative Technologist

Dashverse / Frameo

SUBMITTED BY Manish Das

REFERENCE AD [Chicnutrix "Glow Advanced"](#) · 45.6 s · 9:16 · 17 shots

PART 2 VIDEO [Four Atoms — watch on Google Drive](#)

STACK Claude · Higgsfield MCP · GPT-Image · Premiere Pro + Premiere MCP



What this workflow is for

It is a cloner. What it clones is structure — never a brand, a face, or a claim.

Performance video is a volume problem wearing a creative costume. A brand that finds an ad that works does not need one more beautiful film; it needs fifteen more of *that*, this month, before the format fatigues. Conventional production cannot quote that job. So the winning ad runs until it dies, and the brand starts again from nothing.

Meanwhile the thing worth copying is already public. Every ad still running sits in an open ad library, and what makes it work is not a secret: an eight-second hook before any product appears, a cutting rate that accelerates through the problem and slows at the solution, a three-second logo card. That is architecture. Architecture is not ownable, and it is the part most teams cannot see well enough to reuse.

This workflow measures it, then rebuilds it. It takes a proven ad, extracts the skeleton to the frame, and produces a new one for a different brand with an invented presenter who has never existed. The rhythm is inherited. Everything a viewer could recognise is new.

INHERITED

- Hook, mid, CTA timing
- Cutting rhythm and shot count
- Where B-roll falls
- Order of the argument
- Camera language

INVENTED

- Brand, product, packaging
- The presenter's face
- Script and every claim
- Voice, wardrobe, location
- Music and sound

The hard part was never generating video. It is making one face survive fourteen separate generations without quietly becoming someone else. So the workflow spends its first half locking things down — a character reference sheet, prompt blocks that are never edited, recorded seeds — and its second half checking every output against that lock rather than against the shot before it. Drift cannot accumulate if nothing is ever compared to its neighbour.

The last requirement is that it runs without me. The steps are numbered, the prompts are templates, the rejection criteria are things a non-expert can see by pausing on a frame. A client's own team can produce the next video and the one after that, which is the only version of this that is worth building.

What it is not.

Not a way to pass off another brand's work as your own. The reference supplies timing and nothing else — no logo, no packaging, no script, and no one's likeness. A claim that survives into the new script has to be cleared on its own merits, and the workflow stops and escalates when it isn't.

Contents

Everything in this submission is inside this one file. The workflow files are also embedded as attachments, and the two videos are linked and previewed below — nothing depends on hunting for a separate folder.

PART 1 Production Spec

Teardown of the reference ad: format, camera reasoning, a shot-by-shot table with measured timecodes and stills, locked elements, and three client questions.

PART 2 The Video

The video made for an invented brand, how it maps to the spec, declared deviations, and how character consistency was held.

PART 3 Handoff SOP

Instructions detailed enough for a non-expert to produce the next video in this format without asking me anything. Seven sections.

APPENDIX A ad-cloner.md

The workflow file the process actually runs on, reproduced verbatim.

APPENDIX B prompt_sheet.md

Every prompt as a fill-in-the-blank template, reproduced verbatim.

The two videos



Reference — Chicnutrix “Glow Advanced”

45.6 s · 720×1280 · 30 fps · 17 shots · analysed in Part 1

Watch: [Meta Ad Library](#)



Part 2 — Four Atoms Glutathione Effervescent Tablets

Invented brand · 9:16 · shared anyone-with-link, no sign-in required

Watch: [Google Drive](#)

How the video was made

A competitor ad is selected through an ad-intelligence tool and analysed into a three-part structure — hook, mid-section, CTA — with every cut, B-roll position and camera motion recorded. The script is rewritten for the new brand and, if it mixes Hindi and English, normalised into Devanagari so voice generation pronounces it correctly. An avatar is generated from the reference's age and gender read, then locked to a reference sheet. The script is chunked into beats of no more than ten seconds each, every beat is generated as talking-head video through the Higgsfield MCP, and the clips are placed on a Premiere Pro timeline through the Premiere MCP. The hook is rebuilt separately with motion, the CTA card is composited from the product image, B-roll is generated and trimmed by hand, and SFX and music come from an in-house Premiere library. Full detail in Part 3.

The workflow files are attached to this PDF, not just printed in it.

Both .md files are embedded as file attachments, so the prompts can be extracted and used directly rather than retyped. Open the attachments pane in your PDF reader — the paperclip icon in Acrobat, or *View → Navigation Panels → Attachments*.

The workflow at a glance

One competitor ad in, one finished ad out. Every step below is documented in full in Part 3 — this is the map.

ANALYSE *What is the reference actually doing?*

- 1 Select the competitor ad** Ad Intelligence
Search the brand, filter to long-running ads, download the MP4
- 2 Analyse the video** Claude
Cut detection → shot list, hook / mid / CTA split, B-roll map, hook motion

SCRIPT & CHARACTER *Lock everything that must stay identical*

- 3 Rewrite the script** Claude
Clone-and-swap the product and ingredients, or write a new lead to the same beats
- 4 Normalise Hinglish** Claude
Romanised Hindi → Devanagari, so voice generation says it correctly
- 5 Lock the avatar** QC GATE GPT-Image
Age and gender read from the reference → reference sheet the whole video inherits
- 6 Chunk into beats** Claude
Hook stays whole · every other beat ≤ 10 s · never split a sentence or a Q&A

GENERATE *Make the shots, one beat at a time*

- 7 Generate the A-roll** QC GATE Higgsfield MCP
One talking-head clip per beat, checked against the reference sheet
- 8 Rebuild the hook** QC GATE Higgsfield MCP
Motion prompt from the reference, performed by the locked avatar
- 9 Composite the CTA card** Higgsfield MCP
Product image dropped into the standardised 3–4 s end card
- 10 Generate the B-roll** Higgsfield MCP
Inserts rebuilt at the positions the analysis reported

ASSEMBLE *Cut it together and ship*

- 11 Lay the timeline** Premiere MCP
Clips placed in script order, hook at the head, CTA at the tail
- 12 Trim the B-roll to fit** MANUAL Editor
Inserts overrun their slots — the editor trims by hand, A-roll is never stretched
- 13 Sound, captions, export** Premiere Pro
SFX and BGM from the in-house repo · captions on every frame · H.264 9:16

Figure 1. The full pipeline, start to finish. Four phases, thirteen steps. The chip on each step names the tool that does the work. ■ **QC gate** marks a step where output is checked and rejected before anything downstream is built; ■ **Manual** marks the one step that is deliberately not automated.

PART 1 OF 3

Production Spec

Turning “we want this, but for us” into a written plan another team could execute without ever seeing the original.

Production Spec — Reference Ad Teardown

Reference ad: Chicnutrix "Glow Advanced" — influencer/UGC performance ad **Public link:** [Meta Ad Library — ad ID 1393506182719580](#) **Runtime:** 45.6 s · 720×1280 (9:16 vertical) · 30 fps · 17 shots **Analysed by:** frame-accurate scene-cut detection (threshold 0.25) + shot-by-shot frame extraction

1. What kind of video this is

A 45-second vertical direct-response ad for an ingestible beauty supplement, built as a **single-presenter confessional** — one woman talks straight to camera for the whole runtime while B-roll and product footage cut in over her voice. It opens on a 8-second contrarian hook ("*The worst thing Instagram ever did was convince me that every skin problem needed a different product*") that attacks the category before it sells anything.

The job it does for the brand: **reframe the buyer's problem so the product becomes the obvious answer**. It is not an awareness film. It argues that a shelf full of topical products is the wrong solution, then positions one ingestible as the replacement — and closes on a time-bound proof claim ("in 6 weeks* only") designed to drive a click, not a feeling.

2. How it's shot and why

CHOICE	WHY IT WORKS
Single fixed medium shot, waist-up, eye-level for every A-roll beat	Costs nothing to shoot, and the unchanging frame makes the presenter feel like a person talking to you rather than an ad. Stability is also what lets 6 separate A-roll segments cut together as one continuous conversation.
Same room, same outfit, same lighting for 100% of A-roll	The ad cuts away 10 times. Every return to the presenter must feel like <i>no time has passed</i> — one wardrobe and one setup is what buys that.
Soft warm key from frame-left, visible palm-frond shadow on a beige wall	Reads as "someone's apartment in the afternoon," not a studio. The shadow is doing heavy lifting: it's a fixed, recognisable background feature that proves continuity across every cutback.
Extreme close-up B-roll of real skin texture (pigmentation, breakouts, pores)	The problem has to look real and unretouched, or the promise is not credible. These are the only shots in the ad that are deliberately unflattering.
Fast cutting through the problem list — 4 shots in 4.2 seconds (6.7 s → 10.8 s)	The cutting rate <i>is</i> the argument. Each rapid cut is one more product you supposedly need; the edit makes the category feel exhausting before the script says so.
Cutting slows to 3-4.7 s per shot after the product appears	Pace drop signals relief. The problem section is frantic; the solution section breathes.
CG liposome renders (28.9 s → 36.8 s, 7.9 s total — the longest single block)	Supplements can't show results in action, so abstract science imagery substitutes for proof. It also buys the ad its most "premium" visual without a production budget.

CHOICE	WHY IT WORKS
Kinetic caption on every single frame, mixed serif-italic and sans	Built for sound-off autoplay. The italic serif carries the emotional words, the sans carries the claims — so the hierarchy survives even as a thumbnail.
Before/after split at 39.2 s, held only 1.7 s	Placed late, after credibility is built, and held short — long enough to register, short enough not to invite scrutiny.
3.4-second logo-only end card	Standard performance-ad practice: brand last, once the argument has landed.

3. Shot-by-shot breakdown

Structure: Hook 00.00–06.10 (8.1 s) · Mid 8.10–42.23 (34.1 s) · CTA 42.23–45.59 (3.4 s) **Shot mix:** 6 A-roll segments · 10 B-roll/CG inserts · 1 end card

#	FRAME	IN - OUT	TYPE	WHAT'S IN FRAME	WHAT HAPPENS
1		00.00–06.67 6.67s	A-ROLL	MS presenter, pink dress, beige wall, white facial roller	Full hook line delivered to camera; gestures with roller
2		06.67–08.10 1.43s	B-ROLL	ECU cheek, visible pigmentation	“a cream for pigmentation”
3		08.10–09.77 1.67s	B-ROLL	ECU neck / jawline	“a serum for hydration”
4		09.77–09.90 0.13s	A-ROLL	MS presenter	Single-frame flash cutback — rhythmic punctuation
5		09.90–10.83 0.93s	B-ROLL	ECU under-eye, finger touching skin	“a treatment for breakouts”
6		10.83–12.43 1.60s	B-ROLL	CU green clay mask, gua-sha tool	“a mask for dullness”
7		12.43–15.77 3.34s	A-ROLL	MS presenter	“and somehow...” — pace resets, turn into the pitch
8		15.77–18.80 3.03s	PRODUCT	3 tubes on plinth, brand lockup	Product reveal, name on screen
9		18.80–23.07 4.27s	B-ROLL	Effervescent tablet dropping into water glass	Hero ingredient named on screen
10		23.07–25.30 2.23s	B-ROLL	CU presenter drinking through straw	“antioxidant” — shows the ritual

#	FRAME	IN - OUT	TYPE	WHAT'S IN FRAME	WHAT HAPPENS
11		25.30–28.90 3.60s	A-ROLL	MS presenter	"to slow signs of..."
12		28.90–33.57 4.67s	CG	Liposome spheres, iridescent	Delivery-technology claim
13		33.57–36.83 3.26s	CG	Bubbles inside membrane	"for faster absorption and impact"
14		36.83–39.17 2.34s	A-ROLL	MS presenter	"11 more skin-lovin' ingredients"
15		39.17–40.83 1.66s	B-ROLL	Before/after split, same face	"Get visible results"
16		40.83–42.23 1.40s	A-ROLL	MS presenter	"In 6 weeks* only"
17		42.23–45.59 3.36s	CTA	Logo end card on cream	Brand lockup + tagline

■ Hook ■ Mid-section ■ CTA · Timecodes measured by scene-cut detection at threshold 0.25 on the source file.

4. What must never change

These are locked. Any future video in this series that breaks one of them stops feeling like part of the same family.

Character - One presenter, female, mid-20s, South Asian, shoulder-length dark hair worn down - Same face across every A-roll shot — this is the single hardest constraint and the first thing a viewer notices when it breaks - One outfit per video, saturated single-colour, spaghetti-strap — never patterned - Delivers to lens, never to an off-camera interviewer

Environment - One room per video, shot from one camera position for all A-roll - Warm beige / sand wall, visibly textured, never pure white - Hard plant-shadow cast frame-left — the continuity anchor - Nothing in frame changes position between cutbacks

Look - Warm palette only: beige, sand, terracotta, skin tones, plus one saturated accent (here, magenta) carried by wardrobe - Soft directional key from frame-left, no hard shadow on the face, one lighting logic for the whole video - 9:16, 720×1280 minimum, 30 fps

Structure - Hook 8–10 s, single unbroken A-roll, no cut inside it - CTA 3–4 s, logo-only, no presenter - Mid-section beats never exceed 10 s - Captions on 100% of frames — the ad must work muted

Tone - First person, conversational, mild contrarian irritation — never bubbly, never clinical - The problem is named before the product appears - One time-bound claim near the end, asterisked

5. Three questions for the client

1. **Where does this run — Meta Reels, YouTube pre-roll, or Amazon PDP?** Reels can carry an 8-second hook before the payoff; a skippable YouTube pre-roll needs the product visible before 0:05, which would restructure the entire first third.
2. **What claims has your regulatory/legal team already cleared, and what does the "6 weeks*" asterisk resolve to?**
The whole ad is built on a time-bound efficacy claim. If that number isn't substantiated for this market, the script's spine changes, not just its wording.
3. **Is this presenter a contracted brand face for the series, or one of a rotating cast?** A locked face means we build a reusable character reference set and every future video inherits it; a rotating cast means consistency lives in the room and the grade instead, which is a different build.

PART 2 OF 3

The Video

The proof. Same format, same camera language, invented brand.

Part 2 — The Video

Brand: Four Atoms (*invented for this assignment — not a real company*) **Product:** Four Atoms Glutathione Effervescent Tablets **Runtime:** [XX] seconds · 9:16 vertical · [N] shots **Character appears in:** [N] of [N] shots

Watch it

► [Watch the video](#)

`https://drive.google.com/file/d/12d0oep34td6RuWrTstSrx0syWa9Id_C/view`

Sharing is set to *anyone with the link can view* — no sign-in required. The file is also in this folder as `02_Video/final_video.mp4` .

The product



Four Atoms — L-Glutathione, Radiance + Glow

15 effervescent tablets · watermelon flavour

The pack shot used for the CTA card. It meets the input requirement in the SOP's checklist — clean plain background, product upright and centred, label legible at thumbnail size — which is what lets the standardised CTA system prompt composite it without a retouching pass.

Invented brand and packaging, created for this assignment.

How it maps to the spec

SPEC REQUIREMENT (PART 1)	WHAT THE VIDEO DOES
Hook 8-10 s, single unbroken A-roll	[FILL IN]
Mid-section beats $\leq$ 10 s	[FILL IN]
CTA 3-4 s, logo only, no presenter	[FILL IN]
One room, one camera position, one wardrobe	[FILL IN]
Warm palette + one saturated accent	[FILL IN]
Captions on every frame	[FILL IN]

Where I deviated from my own spec, and why

Part 2 allows deviation as long as it is declared. List each one. If there were none, say so explicitly — an empty section reads as an oversight.

1. **[WHAT CHANGED]** — [WHY. e.g. "The reference runs 45 s; the brief caps this at 30 s, so the mid-section carries 4 beats instead of 8. The hook and CTA windows are unchanged."]

2. [WHAT CHANGED] — [WHY]

Consistency notes

The single most heavily weighted criterion is that the character stays the same person across shots. How that was held here:

- Avatar locked to a fixed reference sheet before any clip was generated; every clip was checked against that sheet rather than against the previous clip, so drift could not accumulate.
- One room, one camera position and one wardrobe for all A-roll.
- Seeds recorded for every approved generation, so nothing needed re-rolling.
- Inserts kept below the nose line — no second face enters the video.

Originality

The brand, product name, packaging and presenter are all invented. No element of the reference ad's identity appears: not its name, logo, packaging, or the likeness of the person in it. What carries over is structure only — hook mechanics, cutting rhythm, and the order in which the argument is made.

PART 3 OF 3

Handoff SOP

Making myself unnecessary — the method as a system someone else runs.

SOP — Producing the Next Video in This Format

Audience: an operator on the client's team who is smart but is *not* an AI-video expert. **Goal:** you can produce a complete, on-format video without messaging anyone. **Workflow file attached:** `workflow/ad-cloner.md` (the end-to-end runbook Claude executes) and `workflow/prompt_sheet.md` (all prompts, copy-pasteable).

1. Format overview

This format is a **45–60 second vertical talking-head ad** for a consumer product. One presenter — an AI-generated character we call the *avatar* — speaks to camera for the entire video from a single fixed medium shot in a single room. Over her voice, short inserts cut in: close-ups of the problem, the product, the ritual of using it, and a before/after. She is the only person who ever appears.

Every correct video in this format has the same three-part spine:

- **Hook — 8 to 10 seconds.** One unbroken shot of the avatar. No cuts inside it. This is where the video either keeps the viewer or loses them.
- **Mid-section — the bulk of the runtime.** Broken into *beats* of **no more than 10 seconds each**, alternating avatar shots with inserts.
- **CTA — 3 to 4 seconds.** A product/logo card. The avatar does not appear.

A video is **correct** when: the avatar's face is recognisably the same person in every shot she appears in; the room, wardrobe and lighting never change within a video; captions are on every frame; and the hook and CTA fall inside their time windows. A rough-looking video that holds all of these is correct. A beautiful one where the face drifts between shots is not.

ANALYSE *What is the reference actually doing?*

- 1 **Select the competitor ad** Ad Intelligence
Search the brand, filter to long-running ads, download the MP4
 - 2 **Analyse the video** Claude
Cut detection → shot list, hook / mid / CTA split, B-roll map, hook motion
- ## SCRIPT & CHARACTER *Lock everything that must stay identical*
- 3 **Rewrite the script** Claude
Clone-and-swap the product and ingredients, or write a new lead to the same beats
 - 4 **Normalise Hinglish** Claude
Romanised Hindi → Devanagari, so voice generation says it correctly
 - 5 **Lock the avatar** QC GATE GPT-Image
Age and gender read from the reference → reference sheet the whole video inherits
 - 6 **Chunk into beats** Claude
Hook stays whole · every other beat ≤ 10 s · never split a sentence or a Q&A

GENERATE *Make the shots, one beat at a time*

- 7 **Generate the A-roll** QC GATE Higgsfield MCP
One talking-head clip per beat, checked against the reference sheet
- 8 **Rebuild the hook** QC GATE Higgsfield MCP
Motion prompt from the reference, performed by the locked avatar
- 9 **Composite the CTA card** Higgsfield MCP
Product image dropped into the standardised 3–4 s end card
- 10 **Generate the B-roll** Higgsfield MCP
Inserts rebuilt at the positions the analysis reported

ASSEMBLE *Cut it together and ship*

- 11 **Lay the timeline** Premiere MCP
Clips placed in script order, hook at the head, CTA at the tail
- 12 **Trim the B-roll to fit** MANUAL Editor
Inserts overrun their slots — the editor trims by hand, A-roll is never stretched
- 13 **Sound, captions, export** Premiere Pro
SFX and BGM from the in-house repo · captions on every frame · H.264 9:16

Figure 1. The full pipeline, start to finish. Four phases, thirteen steps. The chip on each step names the tool that does the work. ■ **QC gate** marks a step where output is checked and rejected before anything downstream is built; ■ **Manual** marks the one step that is deliberately not automated.

2. Before you start: inputs checklist

Do not begin until every box is ticked. Missing item 3 or 4 will cost you the most rework.

Creative inputs - ☐ **Reference ad** — the competitor video you are rebuilding against, downloaded as MP4. Sourced from the Ad Intelligence tool (Meta Ad Library, Foreplay, or equivalent): search the competitor, filter to videos running longest, download one. - ☐ **Script** — either (a) the reference script with product name and ingredients swapped, or (b) a new script you wrote that follows the reference's beat flow. Decide which before you start. - ☐ **Product images** — clean pack shot on transparent or plain background, minimum 1000 px on the long edge. Needed for the CTA card. - ☐ **Brand assets** — logo (PNG, transparent), brand colours (hex codes), and the approved tagline. - ☐ **Claim clearances** — written confirmation of any efficacy or time-bound claim in the script ("in 6 weeks"). Do not generate a video around an uncleared claim.

Character inputs - ☐ **Avatar reference set** — either the locked reference sheet from a previous video (preferred — reuse guarantees consistency), or nothing, if you are creating a new avatar in Step 2.

Tools and access - ☐ **Claude** with the workflow file `workflow/ad-cloner.md` loaded - ☐ **Higgsfield MCP** connected — handles all image and video generation. Confirm credits are available before starting; a mid-run credit failure loses the generation. - ☐ **GPT-Image** access — used for avatar reference generation - ☐ **Adobe Premiere Pro** with the **Premiere MCP** connected, and a project open - ☐ **SFX + BGM repo** installed — our curated Premiere library of transitions, whooshes and beds for this format

3. Step-by-step run instructions

Stage A — Analyse the reference (≈10 min)

1. **Load the reference MP4** into the workflow and run the analysis step. It returns: total runtime, every scene-cut timestamp, and the shot list split into hook / mid / CTA.
2. **Read the B-roll report.** It tells you how many inserts the reference uses and where each one falls. You will rebuild against these positions, so check the numbers look sane before continuing — if it reports 2 shots in a 45-second ad, the cut detection failed; re-run it.
3. **Confirm the three-part split.** Hook should land at 8–10 s and CTA at 3–4 s. If the reference doesn't split that way, it is a poor fit for this format — pick a different reference.

Stage B — Script (≈15 min)

4. **Choose your script mode:** - **Clone mode** — the workflow takes the reference script and replaces only the product name and ingredient names. Fastest, and the beat timing is guaranteed to fit. - **Manual lead mode** — you supply your own opening, and the workflow fits it to the reference's flow. Use this when the client wants a different angle.
5. **Run Hinglish normalisation.** If the script mixes Hindi and English, this step rewrites Hindi words from Roman letters into Devanagari (`kiya` → `किया`). **Do not skip this.** Voice generation mispronounces Romanised Hindi; Devanagari fixes the phonetics.
6. **Read the normalised script aloud against a timer.** Hook must land in 8–10 seconds at natural speaking pace. If it runs long, cut words now — not after you have generated video.

Stage C — Avatar (≈20 min)

7. **Generate or import the avatar.** - *Generate:* the workflow reads the age and gender of the presenter in the reference and returns 2–3 avatar prompt variants. Generate each with GPT-Image and pick one. - *Import:* supply your own character images instead. **Always prefer this if a previous video in the series exists** — reusing the locked reference set is the single most reliable way to keep the face consistent.

8. **Build the reference sheet.** Save 3–4 images of your chosen avatar: front, three-quarter left, three-quarter right, and one at the framing you'll shoot. This sheet is the ground truth for every quality check that follows. Save it to the project folder.
9. **Run the gender/grammar QC.** If the script is in Hindi or Hinglish, the workflow checks that verb gender agrees with the avatar's gender — करा था for a male presenter, करी थी for a female one. Mismatched gender agreement is the most common error in Hindi scripts and is immediately obvious to a native speaker.

Stage D — Chunk the script (≈10 min)

10. **Split the script into generation chunks.** The workflow does this automatically under three hard rules:
 - **Hook** = one chunk, 8–10 s, never subdivided.
 - **CTA** = handled separately in Stage G, not chunked here.
 - **Mid-section** = beats of **maximum 10 seconds**. A question and its answer must stay in the same chunk. Never end a chunk mid-sentence.
11. **Eyeball the chunk list.** If any beat exceeds 10 s or ends on a comma, fix it by hand before generating. Fixing it here costs a minute; fixing it after generation costs a full regeneration.

Stage E — Generate and place the A-roll (≈45 min, mostly waiting)

12. **Generate every mid-section beat** as an end-to-end talking-head video via Higgsfield MCP, using the avatar reference sheet plus motion direction pulled from the reference video's analysis. All A-roll is talking head — one framing, one room.
13. **Quality-check each clip as it lands** against the checks in Section 5. Reject and regenerate before moving on. Do not batch the checking to the end.
14. **Push the approved A-roll to Premiere.** The Premiere MCP lays the clips onto the timeline in script order. You should now have a complete, gap-free spoken track with no inserts yet.

Stage F — Rebuild the hook (≈20 min)

15. **Generate the hook separately, with motion.** The hook is not a static talking head — in most references it carries a distinct camera angle or a physical action (a turn, a gesture, a reach for the product) that is doing the stopping. The workflow extracts that motion from the reference analysis and writes a motion prompt to recreate it with *your* finalised avatar.
16. **Generate 2–3 hook variants and pick the strongest.** This is the highest-leverage 10 seconds in the video; it is worth the extra generations.
17. **Place the hook at the head of the timeline** via the Premiere MCP.

Stage G — CTA (≈10 min)

18. **Supply the product image.** The workflow applies a standardised CTA system prompt that composites your product into the locked CTA layout.
19. **Place it at the tail of the timeline, trimmed to 3–4 seconds.** Not 2, not 6.

Stage H — B-roll and finish (≈45 min, manual)

20. **Generate the inserts** the reference used, at the positions the analysis reported.

21. **Place them on the timeline — this step is manual.** Generated B-roll frequently runs longer than the gap it needs to fill. Trim each insert to fit; do not stretch the A-roll to accommodate it. **This is editor's work and is expected to need hands.**
22. **Add SFX from the repo.** Transitions on every cut into and out of a B-roll insert. Keep it light.
23. **Add BGM from the repo.** Pick one bed for the whole video. Duck it under the voice. You may also import your own track if nothing in the repo fits.
24. **Add captions to every frame.** Sound-off is the default viewing condition; a video without captions is not finished.
25. **Export** — H.264, 9:16, 1080×1920, 30 fps.

4. Prompt templates

Copy-pasteable versions of all of these are in `workflow/prompt_sheet.md`.

How to read these: `[SQUARE BRACKETS]` = you change this for every video. **Bold text** = locked, never edit — this is what keeps the character and style consistent.

4.1 Avatar reference generation

A [AGE RANGE, e.g. 24-28] year old [GENDER] [ETHNICITY] person,
[HAIR DESCRIPTION], [WARDROBE: single saturated colour, no pattern],
**photographed at eye level, medium shot from the waist up,
soft warm directional key light from frame left,
plain warm beige textured wall background with a soft plant-frond shadow,
natural skin texture with visible pores, no retouching,
shallow depth of field, 9:16 vertical, photorealistic, shot on 35mm**

Locked because: the framing, light direction, wall and shadow are the continuity anchors. Change any of them and cutbacks stop matching.

4.2 A-roll beat (talking head)

Reference image: [ATTACH AVATAR REFERENCE SHEET]
The character speaks this line to camera: "[SCRIPT LINE FOR THIS BEAT]"
Expression: [NEUTRAL / CONCERNED / REASSURING / AMUSED]
**Same character as reference image, identical face, identical hair,
identical wardrobe. Static camera, eye level, medium shot waist up.
Warm beige wall with plant shadow, unchanged. Soft key from frame left.
Natural conversational hand gestures. No camera movement. No zoom.
No cuts. Single continuous take.**

Locked because: "no camera movement, no cuts" is what lets six separately generated clips read as one conversation.

4.3 Hook (motion)

Reference image: [ATTACH AVATAR REFERENCE SHEET]
The character [MOTION EXTRACTED FROM REFERENCE – e.g. "turns from three-quarter to face camera and raises the product into frame"] while saying: "[HOOK LINE, 8-10 SECONDS]"
Same character as reference image, identical face, hair and wardrobe. Same warm beige room, same plant shadow, same soft key from frame left. One continuous take, no cuts. Motion completes within the first 3 seconds, then settles to a stable medium shot.

Locked because: the motion must resolve early — a hook still moving at second 7 has no stable frame to cut out of.

4.4 B-roll insert

[SUBJECT: e.g. "extreme close-up of skin texture on a cheek, visible pigmentation"]
Macro detail shot, no face visible above the nose, warm natural light matching a beige interior, shallow depth of field, photorealistic, natural unretouched skin, 9:16 vertical. No text, no logos, no hands unless specified.

Locked because: inserts must colour-match the A-roll room or the cut announces itself. Keeping faces out avoids a second, inconsistent face entering the video.

4.5 CTA card

Product image: [ATTACH PRODUCT PNG]
Brand logo: [ATTACH LOGO PNG]
Tagline: "[TAGLINE]"
Product centred on a plain warm cream background, soft even studio light, subtle contact shadow beneath the product, logo lockup centred above or below the product, generous margins, no clutter, no people, 9:16 vertical.

5. Quality checks

Run the check for a stage **before** moving to the next one. Every check below is something you can verify by looking — no expertise required.

After avatar generation

- ☐ **Open the reference sheet images side by side.** If you would not say these are photos of the same person on different days, regenerate.
- ☐ Hands visible? Count the fingers. Five per hand, or regenerate.
- ☐ Eyes both pointing the same direction.

After each A-roll clip

- ☐ **Pause on the first frame and the last frame. Put each next to the reference sheet. If the face doesn't match, reject and regenerate** — do not tell yourself it will be fine at speed. It will not.
- ☐ **Wall and shadow in the same position** as in the previous approved clip.
- ☐ **Same wardrobe** — check the neckline and any jewellery specifically; these drift first.
- ☐ Mouth movement matches the words; no mouth still moving after the line ends.
- ☐ Camera has not drifted, zoomed or wobbled.
- ☐ No extra limb, no hand passing through the body.

After the hook

- ☐ Time it. **8 to 10 seconds. Outside that range, regenerate.**
- ☐ The motion finishes in the first 3 seconds and the frame is stable after it.
- ☐ Same face as the reference sheet — check this again even though you checked the A-roll.

After B-roll generation

- ☐ **Does the light come from the same side as in the A-roll?** If the A-roll is lit from the left and the insert is lit from the right, reject.
- ☐ Warm tone matches — hold it next to an A-roll frame. A cool-toned insert breaks the room.
- ☐ No unintended faces, text, logos or watermarks.

After the CTA card

- ☐ Product is the correct product, right way up, label readable.
- ☐ Logo is the supplied file, not a regenerated imitation. **A regenerated logo is always wrong — replace it with the real PNG.**
- ☐ Card duration is 3–4 seconds.

Before export

- ☐ **Scrub the whole video at normal speed with sound off.** Does it make sense from captions alone? If not, fix the captions.
- ☐ **Scrub it again watching only the face.** Any shot where she becomes a different person, replace that shot.
- ☐ No beat longer than 10 seconds.
- ☐ Audio has no gaps or overlaps between beats.
- ☐ BGM sits under the voice, never over it.
- ☐ Total runtime inside the target for the placement.

6. Stopping and resuming

You can stop after any stage. Before you walk away, save these — this is the whole list:

Always save - The **Premiere project file** (`.prproj`) — save, don't just leave it open - The **avatar reference sheet** images, in the project folder, named `avatar_ref_01.png` onwards - The **seed or generation ID** of every approved clip,

in a plain text file next to the project. *Without seeds you cannot reproduce a look, and a re-generation will give you a different face.* - The **normalised script** with the chunk splits marked - The **analysis output** from Stage A (the shot list and B-roll timings)

Folder layout to save into

```
project_name/
├─ project_name.prproj
├─ script_normalised.md
├─ analysis_output.md
├─ seeds.txt
├─ avatar_ref/
├─ approved_clips/
└─ rejected/          ← keep these; useful when comparing drift
```

To resume: 1. Open the `.prproj`. The timeline holds everything already placed. 2. Open `seeds.txt` and `script_normalised.md` side by side — together they tell you the last chunk that was approved. 3. Restart at the **first chunk with no approved clip**. Do not regenerate approved clips; their seeds are locked and a regeneration will drift the face. 4. Before generating anything new, re-open the avatar reference sheet. Every new clip is checked against that sheet, not against the last clip you made.

7. When to call for help

Stop and escalate. Do not push through any of these.

SITUATION	WHY YOU MUST ESCALATE
The client wants a different character.	The avatar is a locked setting the whole series inherits. A new character means a new reference sheet and a re-run of every A-roll clip — it is a new build, not an edit.
The face will not stay consistent after 3 regeneration attempts.	Something upstream is wrong — usually a weak reference sheet. More attempts will not fix it and will burn credits.
The script contains a health, efficacy or time-bound claim that isn't on the cleared list.	Regulatory risk. Never generate around an uncleared claim, even as a draft.
The client asks to change the hook length, the CTA length, or the 10-second beat rule.	These define the format. Changing them is a format decision, not a production one.
The reference ad doesn't split into hook / mid / CTA.	It's the wrong reference for this system. Choosing a replacement is a creative call.
Higgsfield or the Premiere MCP disconnects mid-run, or credits run out.	Save the project and the seeds immediately, then escalate. Do not start re-generating to "catch up" — you will lose the seeds that keep the face consistent.
The client supplies a real person's likeness to use as the avatar.	Rights and consent issue. Always escalate.

SITUATION	WHY YOU MUST ESCALATE
The brand logo only exists as a low-res image.	A regenerated or upscaled logo is never acceptable on a CTA card. We need the vector from the client.

APPENDIX A

Workflow File

`ad-cloner.md` — the runbook the pipeline executes, verbatim. Also attached to this PDF as a file.

Ad Cloner – End-to-End Workflow Runbook

> This is the workflow file referenced by `03\_SOP.md` §3. It is executed by Claude with the
> Higgsfield MCP (image + video generation) and the Premiere MCP (timeline assembly) connected.
> Every video generation and image generation in this workflow bills to Higgsfield.

Pipeline

...

[1] Competitor discovery	Ad Intelligence tool -> shortlist -> select reference MP4
[2] Video analysis	cut detection -> shot list -> hook / mid / CTA split B-roll count + timestamps + motion description
[3] Script rewrite	mode A: clone + swap product/ingredients mode B: manual lead, fitted to reference beat flow
[4] Hinglish normalisation	Roman Hindi -> Devanagari (phonetics fix)
[5] Avatar	age/gender from [2] -> 2-3 prompt variants -> GPT-Image OR import existing avatar -> gender/grammar QC for Hindi verb agreement
[6] Chunking	hook = 1 chunk (8-10s, atomic) CTA = standardised system prompt, handled at [9] mid = beats, max 10s, sentence- and Q&A-safe
[7] A-roll generation	Higgsfield: talking-head video per beat -> Premiere MCP: place sequentially on timeline
[8] Hook rebuild	motion prompt from [2] + finalised avatar -> Premiere MCP: place at head
[9] CTA	product image + standardised CTA system prompt -> Premiere MCP: place at tail, 3-4s
[10] B-roll	regenerate inserts at positions from [2] -> Premiere MCP: place, MANUAL trim by editor
[11] Finish	SFX from repo -> BGM from repo -> captions -> export

...

Stage contracts

[2] Video analysis – output schema

```
```yaml
runtime_s: float
fps: int
resolution: [w, h]
cuts: [float] # scene-cut timestamps
sections:
 hook: {in: 0.0, out: float} # target 8-10s
 mid: {in: float, out: float}
 cta: {in: float, out: float} # target 3-4s
shots:
 - {n: int, in: float, out: float, dur: float,
 type: a_roll|b_roll|cg|product|cta,
 framing: str, content: str, caption: str}
broll:
 count: int
```

```

 positions: [{in: float, out: float, subject: str}]
hook_motion: str # feeds stage [8]
...

[4] Hinglish normalisation – rule
Romanised Hindi tokens are rewritten to Devanagari; English tokens are left alone.
Rationale: TTS and lip-sync mispronounce Romanised Hindi (`kiya` reads as English).
...
"maine ye 6 hafte kiya" -> "मैंने ये 6 हफ्ते किया"
"my skin barrier" -> "my skin barrier" (unchanged)
...

[5] Avatar – gender/grammar QC
If script language is Hindi or Hinglish, verb gender must agree with avatar gender.
...
avatar.gender == female -> करी थी / की थी (करा था is an error)
avatar.gender == male -> करा था / किया था (करी थी is an error)
...

Fails the stage and returns to [4] on mismatch.

[6] Chunking – hard rules
1. Hook is one chunk. Never subdivided. Target 8-10s at natural pace.
2. Mid-section beat duration <= 10.0s.
3. A chunk never ends mid-sentence.
4. A question and its answer never split across chunks.
5. CTA is excluded – it is a generated card, not a spoken beat.

[7] / [8] A-roll and hook – locked prompt segments
The bold/locked blocks in `prompt_sheet.md` are emitted verbatim on every generation.
They carry: framing, eye level, key-light direction, wall + plant shadow, wardrobe,
"no camera movement", "single continuous take". These are what make independently
generated clips cut together as one scene.

[10] B-roll – known manual step
Generated inserts routinely overrun their slot. The editor trims each insert to the
gap; the A-roll is never stretched to fit. This is the one stage that is not automated
and is expected to need hands.

Tooling

| Stage | Tool |
|---|---|
| 1 | Ad Intelligence tool (Meta Ad Library / Foreplay) |
| 2 | Claude – analysis from the MP4 |
| 3, 4, 6 | Claude |
| 5 | GPT-Image (avatar stills) |
| 7, 8, 9, 10 | Higgsfield MCP – all image and video generation |
| 7, 8, 9, 10 | Premiere MCP – timeline placement |
| 11 | Premiere Pro + in-house SFX/BGM repo |

Credits
All image and video generation is performed through the Higgsfield MCP.

```

## APPENDIX B

# Prompt Sheet

---

`prompt_sheet.md` — every prompt, copy-pasteable. Also attached to this PDF as a file.

## # Prompt Sheet – copy-pasteable

`[SQUARE BRACKETS]` = change for every video.

Everything **\*\*not\*\*** in brackets is locked – do not edit, it is what keeps the character and the room consistent across separately generated clips.

---

### ## 1. Avatar reference generation (GPT-Image)

```text

A [AGE RANGE] year old [GENDER] [ETHNICITY] person, [HAIR DESCRIPTION], wearing [WARDROBE - single saturated colour, no pattern], photographed at eye level, medium shot from the waist up, soft warm directional key light from frame left, plain warm beige textured wall background with a soft plant-frond shadow, natural skin texture with visible pores, no retouching, shallow depth of field, 9:16 vertical, photorealistic, shot on 35mm

```

Generate 4 images: front, three-quarter left, three-quarter right, and one at final framing. These become the reference sheet.

---

### ## 2. A-roll beat (Higgsfield)

```text

Reference image: [ATTACH AVATAR REFERENCE SHEET]

The character speaks this line to camera: "[SCRIPT LINE FOR THIS BEAT]"

Expression: [NEUTRAL / CONCERNED / REASSURING / AMUSED]

Same character as reference image, identical face, identical hair, identical wardrobe. Static camera, eye level, medium shot waist up. Warm beige wall with plant shadow, unchanged. Soft key from frame left. Natural conversational hand gestures. No camera movement. No zoom. No cuts. Single continuous take.

```

---

### ## 3. Hook (Higgsfield – motion)

```text

Reference image: [ATTACH AVATAR REFERENCE SHEET]

The character [MOTION EXTRACTED FROM REFERENCE ANALYSIS] while saying: "[HOOK LINE - must time to 8-10 seconds]"

Same character as reference image, identical face, hair and wardrobe. Same warm beige room, same plant shadow, same soft key from frame left. One continuous take, no cuts. Motion completes within the first 3 seconds, then settles to a stable medium shot.

```

Generate 2-3 variants. Pick the strongest.

---

#### ## 4. B-roll insert (Higgsfield)

```text

[SUBJECT - e.g. "extreme close-up of skin texture on a cheek, visible pigmentation"]

Macro detail shot, no face visible above the nose, warm natural light matching a beige interior, shallow depth of field, photorealistic, natural unretouched skin, 9:16 vertical. No text, no logos, no hands unless specified.

```

---

#### ## 5. CTA card (Higgsfield – standardised system prompt)

```text

Product image: [ATTACH PRODUCT PNG]

Brand logo: [ATTACH LOGO PNG]

Tagline: "[TAGLINE]"

Product centred on a plain warm cream background, soft even studio light, subtle contact shadow beneath the product, logo lockup centred above or below the product, generous margins, no clutter, no people, 9:16 vertical.

```

Trim to 3-4 seconds on the timeline. Replace any generated logo with the real PNG.

---

#### ## 6. Hinglish normalisation instruction

```text

Rewrite every Romanised Hindi word in this script into Devanagari. Leave English words in Latin script. Do not translate. Do not change word order, meaning, or sentence length – this is a script-level transliteration only, to fix TTS phonetics.

Then verify Hindi verb gender agreement against the avatar's gender:

female avatar -> करी थी / की थी

male avatar -> करा था / किया था

Flag any mismatch.

Script:

[PASTE SCRIPT]

```

---

#### ## 7. Chunking instruction

```text

Split this script into generation chunks under these hard rules:

1. The hook is ONE chunk, 8-10 seconds at natural speaking pace, never subdivided.
2. Every other beat is at most 10.0 seconds.
3. No chunk ends mid-sentence.
4. A question and its answer stay in the same chunk.
5. Exclude the CTA – it is a generated card, not a spoken beat.

Return a table: chunk number, text, estimated duration.

Script:

[PASTE NORMALISED SCRIPT]

```